

Supervised Learning Preference Optimization: Rethinking RLHF and DPO as Supervised Learning

Yaoshiang Ho

YAOSHIANG@GMAIL.COM

Abstract

Direct Policy Optimization (DPO) is a popular approach to aligning large language models with human preferences. In this paper, we analyze the underlying math and propose a new algorithm which we call Supervised Learning Preference Optimization (SLPO).

Keywords: Reinforcement Learning from Human Feedback (RLHF), Direct Policy Optimization (DPO)

1 Introduction

Alignment is the task of ensuring that the behavior of a Large Language Model (LLM) is consistent with human preferences.

A key difference between the alignment phase and other phases of training an LLM is that the alignment phase considers full sequences of text, rather than simply predicting the next token, as in the pretraining and supervised fine-tuning (SFT) phases.

The alignment approach popularized by the commercial success of ChatGPT was Reinforcement Learning from Human Feedback (RLHF, Ouyang et al. (2022)). Despite its effectiveness, RLHF requires training a second model, called a reward model, as well as Proximal Policy Optimization (PPO), resulting in a technique that is more complex than the basic supervised learning. It also requires a Kullback-Leibler (KL) divergence term to regularize the changes to the LLM during alignment training.

Direct Policy Optimization (DPO) is a simpler approach to alignment which does not require a secondary reward model. In the paper introducing DPO, the authors examine the underlying approach of RLHF and propose the DPO objective to align the target LLM directly using maximum likelihood estimation (MLE). The key insight from the DPO paper is that an LLM’s outputs can be reparameterized into a reward model using ratios, logs, and the Bradley-Terry model (Bradley and Terry (1952)).

The specific contribution of this paper is to reframe the alignment phase away from reward modeling entirely and treating it simply as a pure supervised learning problem by training a model to align to a directly modified probability distribution. We call this approach Supervised Learning Preference Optimization (SLPO).

2 Related Work

Related work...

3 Preliminaries: DPO

We review and continue the analysis of the DPO objective by its authors.

The DPO objective is defined as follows:

$$L_{\text{DPO}}(\pi_\theta; \pi_{\text{ref}}) = \underbrace{-\mathbb{E}_{(x, y_w, y_l) \sim D} \log}_{1} \left[\underbrace{\sigma}_{2} \left(\underbrace{\beta \log \frac{\pi_\theta(y_w | x)}{\pi_{\text{ref}}(y_w | x)}}_{3} - \underbrace{\beta \log \frac{\pi_\theta(y_l | x)}{\pi_{\text{ref}}(y_l | x)}}_{4} \right) \right]. \quad (1)$$

In the underbraced section 1, we see the standard negative log likelihood (NLL) objective. In the underbraced section 2, we see the Bradley-Terry model ¹. In the underbraced sections 3 and 4, we see how the reference and language model’s predictions are reparameterized into a winning and losing score: they are the log of the ratio of the language model to the reference model, for the winning and losing completion, respectively. This score is later described as the reward function.

With simple algebraic manipulation, we can rewrite this objective as:

$$L_{\text{DPO}}(\pi_\theta; \pi_{\text{ref}}) = \underbrace{-\mathbb{E}_{(x, y_w, y_l) \sim D} \log}_{1} \left[\underbrace{\sigma}_{2} \left(\underbrace{\beta \log \frac{\pi_\theta(y_w | x)}{\pi_\theta(y_l | x)}}_{3} - \underbrace{\beta \log \frac{\pi_{\text{ref}}(y_w | x)}{\pi_{\text{ref}}(y_l | x)}}_{4} \right) \right].$$

We can interpret the undercomponents as follows: The first and second underbraces remain unchanged. The third underbrace is the log of the ratio language models probability of the winner divided by the loser. This ratio is optimized to be bigger, given that the log, sigmoid, and log from underbraces 1, 2, and 3 are all monotonic functions. Since a ratio of probabilities is an odds ratio, we will refer to this as the language model’s log-odds ratio.

The fourth underbrace is reference model’s log-odds ratio. This is a constant per y_w and y_l and is not differentiated. However, these values and their ratio vary across different y_w and y_l - that is, each row of training data will have a different value for this constant. More specifically, since π_θ is initialized as π_{ref} , the difference between underbraces 3 and 4 starts at zero during training, and the shape of the sigmoid function (underbrace 2) and its gradient are known. During training, as the language model’s log-odds ratio increases, the sigmoid function will naturally regularize and decelerate the increase by reducing the magnitude of the gradient (a useful version of the vanishing gradient problem). This is the exact outcome described by the DPO authors in their analysis of the gradient of DPO. *But we have developed a different intuition the DPO loss: rather than reparameterizing the language model’s output into a reward model, we are simply regularizing the optimization of the language model’s log-odds ratio.*

1. A ranking method that is mathematically equivalent and perhaps more widely understood is the ELO score, used to rank Chess players, and, LLMs in the Chatbot Arena (Elo (1978); Chiang et al. (2024)). Both ELO and Bradley-Terry assign scores to players, and pass the difference through a sigmoid function to assess the probability of the LHS player of winning.

Let’s analyze the regularization effect. The DPO authors investigated the gradient of the sequence of tokens... let’s go two steps further by considering each token individually, and the softmax logit behind it ².

The variable y is a sequence of tokens, defined as

$$\pi_{\text{ref}}(y \mid x) = \prod_{t=1}^T \pi_{\text{ref}}(y_t \mid x, y_{<t}), \quad (2)$$

where $y = (y_1, y_2, \dots, y_T)$ is the output sequence, x is the input context, and $y_{<t} = (y_1, y_2, \dots, y_{t-1})$ represents the tokens generated prior to time step t . The term $\pi_{\text{ref}}(y_t \mid x, y_{<t})$ denotes the conditional probability of generating token y_t given the input x and the previously generated tokens $y_{<t}$.

Since both π language model are LLMs, they are activated using the softmax function. The softmax operation can be mathematically cumbersome due to its dependence on all other logits for normalization. For simplicity, we assume that the logits are transformed using a log-softmax function, which computes the logarithm of the softmax probabilities in a numerically stable way. This transformation results in log-probs, which are easier to reason about since they can be exponentiated to recover probabilities. Importantly, we do not lose any generality because log-probs can also be passed through a softmax function to recover probabilities, ensuring equivalent behavior.

Let us establish the term g for the layers of the model up to the softmax activation, namely, the feature extractor and log-softmax normalization:

$$g(y \mid x) = \text{logsoftmax}(f(y \mid x)) \quad (3)$$

such that $\pi(y \mid x) = \exp(g(y \mid x))$. Plugging Equations 2 and 3 into the DPO objective 1,

$$L_{\text{DPO}}(\pi_{\theta}; \pi_{\text{ref}}) = -\mathbb{E}_{(x, y_w, y_l) \sim D} \log \left[\sigma \left(\beta \log \frac{\prod_{t=1}^{T_w} \exp(g_{\theta}(y_{w,t} \mid x, y_{w,<t}))}{\prod_{t=1}^{T_w} \exp(g_{\text{ref}}(y_{w,t} \mid x, y_{w,<t}))} - \beta \log \frac{\prod_{t=1}^{T_l} \exp(g_{\theta}(y_{l,t} \mid x, y_{l,<t}))}{\prod_{t=1}^{T_l} \exp(g_{\text{ref}}(y_{l,t} \mid x, y_{l,<t}))} \right) \right]$$

which simplifies to:

$$L_{\text{DPO}}(\pi_{\theta}; \pi_{\text{ref}}) = -\mathbb{E}_{(x, y_w, y_l) \sim D} \log \left[\sigma \left(\beta \left(\sum_{t=1}^{T_w} [g_{\theta}(y_{w,t} \mid x, y_{w,<t}) - g_{\text{ref}}(y_{w,t} \mid x, y_{w,<t})] - \sum_{t=1}^{T_l} [g_{\theta}(y_{l,t} \mid x, y_{l,<t}) - g_{\text{ref}}(y_{l,t} \mid x, y_{l,<t})] \right) \right) \right] \quad (4)$$

2. Technically, a logit is the log-odds ratio and it is the output of a feature extractor before a sigmoid activation function in binary classification. The term logit is also used for categorical classification, when a feature extractor’s outputs are activated with the softmax function. However, these softmax logits cannot be exponentiated to calculate an odds ratio. They do have an unfortunately property however: they are shift invariant, leading to numerical instability and the need for log-softmax and its use of the log-sum-exp trick to improve numerical stability. In this paper, since we are dealing with both sigmoid and softmax functions, we will refer to the value passed into a softmax function as a softmax logit.

We notice another possibility to rewrite the DPO objective towards something more familiar. Optimizing the sigmoid of a difference $a - b$ through BCE is equivalent to optimizing the softmax of a and b through CCE towards a one-hot encoded target of $\text{softmax}(a) = 1$ and $\text{softmax}(b) = 2$. This allows us to rewrite the DPO objective as follows:

$$L_{\text{DPO}}(\pi_\theta; \pi_{\text{ref}}) = -\mathbb{E}_{(x, y_w, y_l) \sim D} \text{CCE} \left[y_{\text{true}} = \begin{bmatrix} 1 \\ 0 \end{bmatrix}, y_{\text{pred}} = \text{softmax} \left(\begin{bmatrix} \beta S_w \\ \beta S_l \end{bmatrix} \right) \right] \quad (5)$$

where

$$S_w = \sum_{t=1}^{T_w} (g_\theta(y_{w,t} \mid x, y_{w,<t}) - g_{\text{ref}}(y_{w,t} \mid x, y_{w,<t})),$$

$$S_l = \sum_{t=1}^{T_l} (g_\theta(y_{l,t} \mid x, y_{l,<t}) - g_{\text{ref}}(y_{l,t} \mid x, y_{l,<t})).$$

Key Observation #1: Much of the DPO objective can be simplified to the form of a standard classification task. This formulation of the DPO objective provides our first key observation. What started off as a complex reparameterization of the language model’s output into a reward model and a Bradley-Terry model, has been simplified into a standard categorical crossentropy loss applied to the softmax, in Equation 5. However, the input to the softmax is not a single value, it’s a sum of differences of values. That difference, and that sum, are the remaining nuances to resolve.

Key Observation #2: The ratio of probabilities creating a reward reward function has been transformed into a normalization of logits to make equivalent all the gradients at the start of training. The first remaining nuance is the difference between the language model and the reference model’s logits. Since the language model is initialized by the reference model, the difference starts at zero for both S_w and S_l . This means that at the start of training, all the gradients for all the winning sequences examples will be the same as each other (likewise for the losing sequences).

Key Observation #3: blah. The second remaining nuance is that the softmax is applied to the sum of the inputs, where the inputs were originally softmax logits during the $\text{CCE}(\text{softmax}())$ applied during next token prediction pretraining.

What happens when we the input to a softmax is a sum? This is easy to reason about. At first, the gradient to each S_w will be 0.5, and that gradient will be passed through to each value z_t . In turn, as z_t is updated by a gradient, the collective value of S_w will increase far more than the value of z_t . In fact, during the first step, the update will be roughly $T \times$ greater than the update to z_t , ignoring the effects of adaptive optimizers. So the value of y_w will increase quickly... and in doing so, rapidly approach the

is updated du, the collective update to S_w will increase far more than the value of z_t . In fact, the update will approximate , and the gradient will decrease. This is

= XXX
 XXX
 XXX
 XXX

XXX

XXX

The "logit" to the Each softmax logit of the language model, which was originally conceived to predict the probability of a token, has been normalized by the softmax logit of the reference model via a shift (this normalization should not to be confused with the normalization by the log-softmax function, which was purely for mathematical convenience). The sum of these normalized softmax logits is then passed through a standard sigmoid activation and negative log likelihood function.

The starting value for the difference between the language model and the reference model's logits is zero, since the language model is initialized by the reference model, so the gradient passed to each logit on the first step of optimization is -0.5 for the winning token and 0.5 for the losing token. As the winning sequence increases its odds, and the losing sequence decreases its odds, the gradient will decrease, as is well known for the negative log loss applied to a sigmoid activated function (e.g. classic binary crossentropy loss for logistic regression). Although sigmoid is not an additive function, e.g. $\sigma(a + b) \neq \sigma(a) + \sigma(b)$, it is monotonically increasing, e.g. $\sigma(a + b) > \sigma(a)$ where $b > 0$. So as the softmax logit of each winning token increases and each losing token decreases, the magnitude of the gradient applied to all softmax logits decreases.

And yet, the DPO equation involves significant machinery—such as the Bradley-Terry model, the log of ratios, and the concept of reward functions—for what ultimately reduces to a standard negative log-likelihood loss applied to normalized softmax logits. Could this normalization be further simplified, perhaps eliminating the need for these additional constructs and reframing the problem in a more straightforward way?

4 Preliminaries: Probabilistic Classification

A basic task for deep learning has been classification, both binary and categorical. Which of the ten digits is this image(LeCun (1998))? Is this pixel part of an object's mask? Is this movie review's sentiment positive or negative(Pang and Lee (2004))?

Modern LLMs auto-regressively predict a probability distribution for the next token among a finite set of possibilities, conditioned on the previous tokens. Hence, an LLM also solves a classification problem.

Given the prevalence of classification, deep learning often assumes target labels can be described as either ordinal numbers or equivalently one-hot encoded vectors. When target labels are intermediate values between one and zero, they were called "soft targets", treated as an regularization technique or even called dark knowledge (Hinton et al. (2015); Szegedy et al. (2016); Hinton et al. (2014)).

But the two roots of the crossentropy loss functions: KL divergence and Maximum Likelihood Estimation (MLE) applied to the probability distribution function of a Multi-Bouroulli distribution, are not limited to binary and categorical outcomes. For example, the crossentropy loss function can be used directly to condition a model to predict to be true 60% of the time and false 40% of the time, as demonstrated in Figure 4.

The term we will use for the idea of predicting a probability distribution is *probabilistic classification*.




Figure 1: Setting y_{true} to a value other than 0 or 1 is a valid use of the crossentropy loss function and perfectly consistent with both KL divergence and MLE.

5 Supervised Learning Preference Optimization

Armed with intuition of the DPO loss and probabilistic classification, we can now turn to a simpler, supervised learning approach to alignment.

The SLPO objective is defined as follows:

$$L_{\text{SLPO}}(\pi_{\theta}; \pi_{\text{ref}}) = -\mathbb{E}_{(x, y_w, y_l) \sim D} \left(\underbrace{\log(w_w * \pi_{\theta}(y_w | x))}_1 + \underbrace{\log(w_l * \pi_{\theta}(y_l | x))}_2 + \underbrace{\log((1 - w_w - w_l) * \pi_{\theta}(y_o | x))}_3 \right)$$

where $w_w = \alpha * (\pi_{\text{ref}}(y_w | x) + \pi_{\text{ref}}(y_l | x))$, $w_l = (1 - \alpha) * (\pi_{\text{ref}}(y_w | x) + \pi_{\text{ref}}(y_l | x))$.

Intuitively, the goal is to take the probabilities of the winning and losing sequence and reallocate them. If alpha is zero, then all of the probability will be allocated to the winning sequence. Importantly, the sum of the probabilities of the winning and losing tokens will be the same in the language model and the reference model. Obviously, this implies that the probability of all other sequences are optimized to be the same in the language model and the reference model. This is essentially the SLPO’s version of the KL divergence regularization term in traditional RLHF.

Note that there are no negative signs in the SLPO objective. This is because the SLPO objective is a pure supervised learning objective, and the negative sign in what the DPO paper calls the unlikelihood objective is not consistent with MLE, since MLE always applies the negative log loss to probabilities, and probabilities are always positive.

6 Results

Section body

Here is a citation Chow and Liu (1968).

7 Conclusion

Section body

Here is a citation Chow and Liu (1968).

Acknowledgments and Disclosure of Funding

The author thanks Chiara Cerini, Peter Tran, and Sam Wookey for their invaluable reviews of early drafts of this work.

Appendix A.

[appendix]

Appendix B. Equivalence of Cross-Entropy on Sigmoid of Difference and Categorical Cross-Entropy on Softmax of True Class for Two Classes

Logistic regression uses a single logit α , transforming it through the *sigmoid* function to obtain a probability $\sigma(\alpha) = \frac{1}{1+e^{-\alpha}}$. The *binary cross-entropy* (BCE) loss for a label y and predicted probability $\hat{y} = \sigma(\alpha)$ is:

$$\text{BCE}(y, \hat{y}) = -[y \log(\hat{y}) + (1 - y) \log(1 - \hat{y})].$$

Alternatively, single-label multiclass classification of two outcomes uses two softmax logits z_1 and z_2 (one per class) and applies the softmax function to obtain probabilities:

$$p(y = 1 \mid z_1, z_2) = \frac{e^{z_1}}{e^{z_1} + e^{z_2}}, \quad p(y = 2 \mid z_1, z_2) = \frac{e^{z_2}}{e^{z_1} + e^{z_2}}.$$

The associated categorical cross-entropy (CCE) loss is:

$$\text{CCE}(y, \hat{y}) = - \sum_{k=1}^2 y_k \log(\hat{y}_k),$$

where $\hat{y}_1 = \frac{e^{z_1}}{e^{z_1} + e^{z_2}}$ and $\hat{y}_2 = \frac{e^{z_2}}{e^{z_1} + e^{z_2}}$.

If we set

$$\alpha = z_1 - z_2,$$

then

$$\frac{e^{z_1}}{e^{z_1} + e^{z_2}} = \frac{1}{1 + e^{z_2 - z_1}} = \sigma(\alpha).$$

Therefore, the predicted probability of class 1 in the softmax parameterization matches the $\sigma(\alpha)$ in logistic regression. This means that the BCE loss on the sigmoid of a difference of two numbers ($z_1 - z_2$) is equivalent to the CCE loss on the softmax of the first number (z_1) (Goodfellow et al. (2016)).

Remainder omitted in this sample. See <http://www.jmlr.org/papers/> for full paper.

References

- R. A. Bradley and M. E. Terry. Rank analysis of incomplete block designs: I. the method of paired comparisons. *Biometrika*, 39(3/4):324–345, 1952. doi: 10.2307/2334029. URL <https://doi.org/10.2307/2334029>.
- Wei-Lin Chiang, Lianmin Zheng, Ying Sheng, Anastasios Nikolas Angelopoulos, Tianle Li, Dacheng Li, Hao Zhang, Banghua Zhu, Michael Jordan, Joseph E Gonzalez, et al. Chatbot arena: An open platform for evaluating llms by human preference. *arXiv preprint arXiv:2403.04132*, 2024.

- C. K. Chow and C. N. Liu. Approximating discrete probability distributions with dependence trees. *IEEE Transactions on Information Theory*, IT-14(3):462–467, 1968.
- Arpad E. Elo. *The Rating of Chessplayers, Past and Present*. Arco Publishing, New York, 1978. ISBN 9780668047210.
- Ian Goodfellow, Yoshua Bengio, and Aaron Courville. *Deep Learning*. MIT Press, 2016. URL <http://www.deeplearningbook.org>.
- Geoffrey Hinton, Oriol Vinyals, and Jeff Dean. Dark knowledge. Presented at the University of Chicago, 2014. URL <https://www.ttic.edu/dl/dark14.pdf>.
- Geoffrey Hinton, Oriol Vinyals, and Jeff Dean. Distilling the knowledge in a neural network. *arXiv preprint arXiv:1503.02531*, 2015. URL <https://arxiv.org/abs/1503.02531>.
- Yann LeCun. The mnist database of handwritten digits. <http://yann.lecun.com/exdb/mnist/>, 1998.
- Long Ouyang, Jeffrey Wu, Xu Jiang, Diogo Almeida, Carroll Wainwright, Pamela Mishkin, Chong Zhang, Sandhini Agarwal, Katarina Slama, Alex Ray, et al. Training language models to follow instructions with human feedback. *Advances in neural information processing systems*, 35:27730–27744, 2022.
- Bo Pang and Lillian Lee. A sentimental education: Sentiment analysis using subjectivity summarization based on minimum cuts. In *Proceedings of the ACL*, 2004.
- Christian Szegedy, Vincent Vanhoucke, Sergey Ioffe, Jon Shlens, and Zbigniew Wojna. Rethinking the inception architecture for computer vision. In *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition (CVPR)*, pages 2818–2826, 2016.